

Sneha Reddy

512-555-3456 | sneha.reddy@email.com | <https://linkedin.com/in/snehareddy>

PROFESSIONAL SUMMARY

AI/ML Engineer with 5 years of experience building and deploying machine learning models at scale. Expertise in NLP, computer vision, and MLOps. Strong Python background with production-grade experience in model serving and cloud infrastructure.

EDUCATION

Master of Science in Machine Learning — Carnegie Mellon University (GPA: 3.8/4.0) | Dec 2020

Bachelor of Engineering in Computer Science — BITS Pilani (GPA: 8.8/10) | Jun 2018

PROFESSIONAL EXPERIENCE

Senior ML Engineer — Google DeepMind | Jan 2021 – Present

- Developed large language model fine-tuning pipelines using PyTorch and JAX on TPU clusters
- Built model serving infrastructure using TensorFlow Serving, FastAPI, and Kubernetes
- Created data preprocessing pipelines processing 50TB+ datasets using Apache Beam and BigQuery
- Published 2 research papers on efficient transformer architectures at NeurIPS and ICML

ML Engineer Intern — Microsoft Research | May 2020 – Aug 2020

- Developed a text classification model for Azure Cognitive Services achieving 94% F1 score
- Built evaluation pipelines using Python, MLflow, and Azure ML

PROJECTS

MedScan AI: An AI diagnostic tool for medical imaging using CNNs and transfer learning, deployed on AWS SageMaker with a Flask-based web interface.

SentimentStream: Real-time Twitter sentiment analysis pipeline using Kafka, Spark Streaming, and BERT for financial market prediction.

TECHNICAL SKILLS

Programming Languages: Python, C++, SQL, R, Bash

AI/ML Frameworks: PyTorch, TensorFlow, JAX, Hugging Face Transformers, scikit-learn

Cloud & MLOps: GCP (BigQuery, Vertex AI), AWS (SageMaker, Lambda), Docker, Kubernetes, MLflow

Data Engineering: Apache Beam, Spark, Kafka, Airflow

Backend: FastAPI, Flask (basic web development)

Databases: BigQuery, PostgreSQL, Redis

Version Control: Git, GitHub

Soft Skills: Research, Technical Writing, Mentoring, Cross-functional Communication